

REAL TIME HUMAN DETECTION & TRACKING VIA A QUADCOPTOR AND SSD- MOBILENET NEURAL NETWORK

**Daksh Shukla, Ben Liu, Patrick McNamee,
Tim Fox, Caio Vigo**

OUTLINE

- Background & motivation
- System architecture – hardware
- System architecture – software
- ROS architecture – 2 packages communication
- Neural network architecture
- Controller
- Challenges – implementation (frequency, TF+ROS, Oscillations)
- Non drone videos – lab testing
- Real time drone video – outdoor flying
- Limitations & future work
- References

BACKGROUND AND MOTIVATION

WHAT WE WOULD LIKE TO DO?

We want to be able to **track** and **follow a human being** using low cost drone for surveillance applications.

WHAT IS THE CHALLENGE?

We want to do it in real-time.

We need to deal with **real-time computation costs** related to **vision processing**.

- Classify a human from real-time drone images
- As the person moves, send commands to the drone to follow the classified object (human)
- Algorithms (neural networks) need to interact with other operating systems (drone, ROS, etc.)

HOW ML CAN HELP?

Neural networks are well-suited for object detection and classification within images.

DEEP NEURAL NETWORKS (DNNs)

Due to their inherent parallel processing and robust object detection characteristics provide a good framework for this type of application.

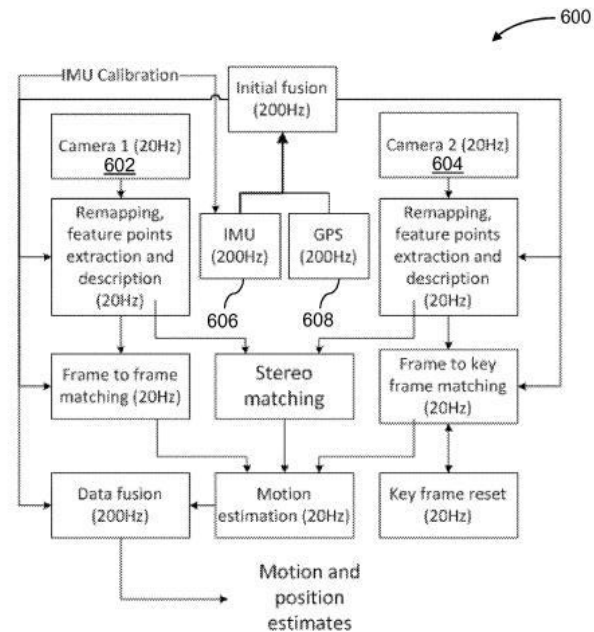


WHO IS CURRENTLY DOING SIMILAR THINGS?

- DJI and others use two cameras for stereo matching [1]
 - Part of obstacle avoidance and position tracking
 - Doesn't classify objects




Cameras for stereo matching



WHO IS CURRENTLY DOING SIMILAR THINGS?

- Researchers have used SLAM (**s**imultaneous **l**ocalization and **m**apping) techniques to recognize objects (no classification) in the environment to follow humans [1].
- Norman Di Palo has a good guide to human tracking using tensor for a quadcopter [2]. Unsure if he has implemented it.



[1] K. K. Lekkala and V. K. Mittal, "Simultaneous aerial vehicle localization and Human tracking," *2016 IEEE Region 10 Conference (TENCON)*, Singapore, 2016, pp. 379-383. doi: 10.1109/TENCON.2016.7848025

[2] <https://medium.com/nanonets/how-i-built-a-self-flying-drone-to-track-people-in-under-50-lines-of-code-7485de7f828e>

SYSTEM
ARCHITECTURE –
HARDWARE



PARROT AR DRONE 2.0

Characteristics:

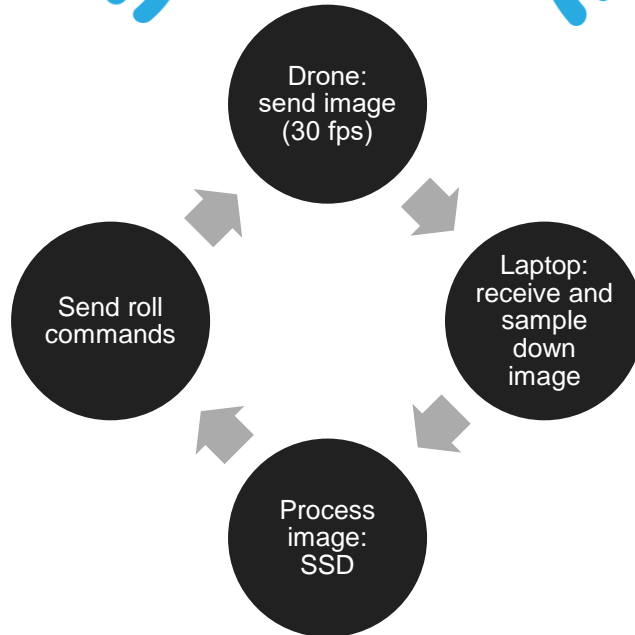
- Front camera: 720p
- WiFi connection: acts as router
- Inertial measurement unit:
Gyroscope, Accelerometer,
Magnetometer
- Altitude ultrasound sensor
- Vertical (bottom) camera
- Take-off weight: 0.93 pounds



SYSTEM ARCHITECTURE - HARDWARE



Parrot AR DRONE 2.0

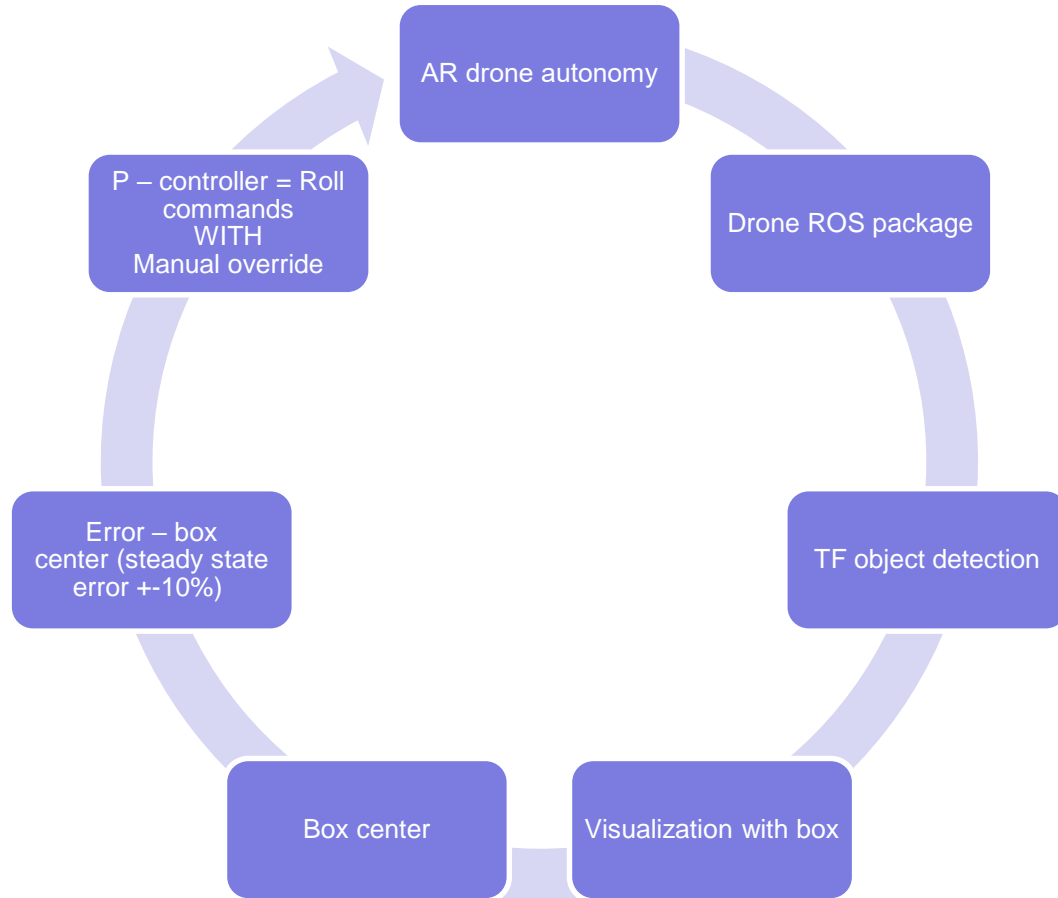


SYSTEM ARCHITECTURE – SOFTWARE

AN OVERVIEW OF THE PROCESS

- Robot Operating System (ROS) software on Ubuntu 16.04
- ROS Kinetic Kame version
- Ardrone autonomy – ROS driver for the drone
 - > *Autonomy lab – Simon Fraser University*
 - > Publishes navigation data and sends commands
- Ardrone tutorials ROS package
 - <https://robohub.org/up-and-flying-with-the-ar-drone-and-ros-getting-started/>
 - > Uses Ardrone autonomy driver to subscribe to sensor data
- Tensorflow object detection API
 - https://github.com/tensorflow/models/tree/master/research/object_detection

SYSTEM ARCHITECTURE - SOFTWARE



ROS ARCHITECTURE - 2 PACKAGES COMMUNICATION

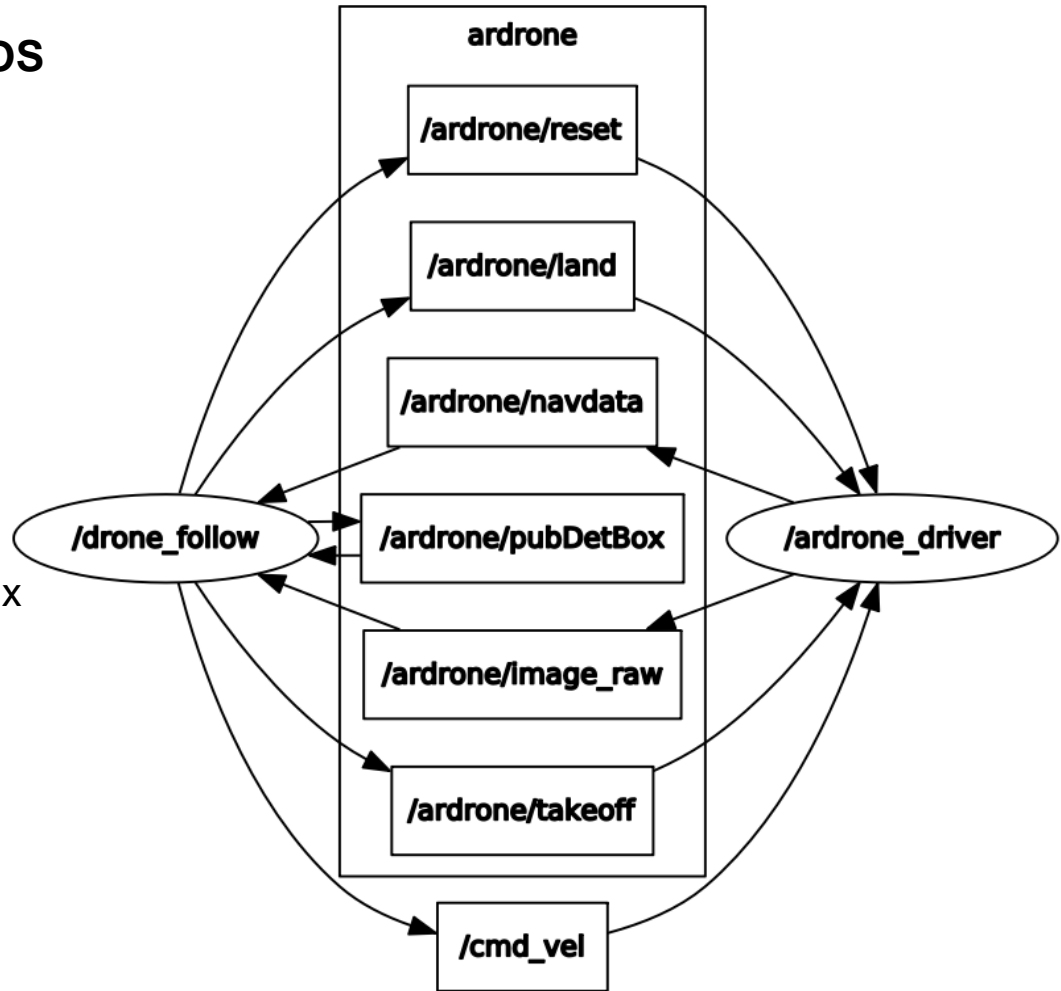
ROS rqt graph:

Visualization of ROS nodes communication



ROS rqt graph: Visualization of ROS topics communication – data exchange

1. Subscribe to image
2. Pass every "nth" (6th) image
frequency = $30/n$ (= 5 Hz)
3. Convert ROS image to CV image
4. Detect classes = ANN(image)
5. If class == 1 && det_score > 30%
find max score index
6. Use index
 - get box = [y_min, x_min, y_max, x_max]
7.
 - visualize(image + box)
 - Make ROS message
 - Publish ROS topic
8. Subscribe to box data
9. Compute error w.r.t center
10. Compute & Publish roll commads



NEURAL NETWORKS ARCHITECTURE

NET ARCHITECTURE

- We used the pre-trained model from TensorFlow object detection API, `ssd_mobilenet_v1_coco_2017_11_17`
- 3,191,072 parameters
- It was pretrained on Microsoft COCO (Common objects in context) data set (available at <http://cocodataset.org/#home>)
Training images = 200,000
Validation images = 8000
- Training algorithm: RMSProp (Root Mean Square Propagation)
batch size = 32
learning rate = 0.004
learning rate decay = 0.95 every 800k steps

J. Huang, V. Rathod, C. Sun, M. Zhu, A. Korattikara, A. Fathi, I. Fischer, Z. Wojna, Y. Song, S. Guadarrama et al., *Speed/accuracy trade-offs for modern convolutional object detectors*, 2016
T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Ramanan, P. Dollár, and C. Lawrence Zitnick. Microsoft COCO: Common objects in context. In *ECCV*, 1 May 2014.1, 4

NET ARCHITECTURE-SSD


- **SSD** stands for **Single Shot MultiBox Detector**
- **Feed forward network:** image through one-time step
- Convolutional feature layers at the end of the base network and they get smaller as the layers move forward, this allows for predictions at different scales
- Default bounding boxes at each feature map location, the SSD will make 4 box predictions and keep the highest score
- SSD will name a 'positive match' if the intersection over union is above 50 percent
- SSD is generally fast and accurate
58 FPS with mAP 72.1% on VOC2007 test
Faster R-CNN: 7 FPS with mAP 73.2%
YOLO: 45 FPS with mAP 63.4%

NET ARCHITECTURE-SSD

- SSD creates **many predictions** for all the objects, there are many negative scores also and the model can hurt from too many negative scores
- SSD sorts negative scores, keeps only 3 negative scores for every positive score
--> keeps the negative scores so that the model can know what a 'bad prediction' is
- SSD uses **non-maximum suppression** to avoid duplicate predictions on objects
- Loss function is as follows where N is the number of positive match and α is the weight for the localization loss

$$L(x, c, l, g) = \frac{1}{N} (L_{conf}(x, c) + \alpha L_{loc}(x, l, g))$$


NET ARCHITECTURE-SSD




NET ARCHITECTURE-MOBILENETS

- **MobileNets** are specialized convolutional neural networks for mobile and embedded application
- **MobileNet** model is based on depthwise convolution and 1×1 convolution called a pointwise convolution.
- The depthwise convolution applies a single filter to each input channel, whereas the pointwise convolution then applies a 1×1 convolution to combine the outputs the depthwise convolution


NET ARCHITECTURE-MOBILENETS



(a) Standard Convolution Filters



(b) Depthwise Convolutional Filters



(c) 1×1 Convolutional Filters called Pointwise Convolution in the context of Depthwise Separable Convolution

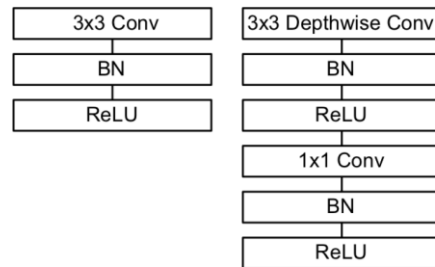


Figure 3. Left: Standard convolutional layer with batchnorm and ReLU. Right: Depthwise Separable convolutions with Depthwise and Pointwise layers followed by batchnorm and ReLU.


M : input channels,
 N : output channels,
 $D_K \times D_K$: the kernel size

CONTROLLER


CONTROL ALGORITHM

- Proportional controller:
Output = $K \cdot \text{Error}$
- Define error:
 $C = (x_{\min} + x_{\max})/2$
If $C > 60\%$ or $C < 40\%$:
 error = $(C - 50\%)$
Else:
 error = 0
roll = $K_r \cdot \text{error}$
- The control law ensures that the person is laterally centered in the image with $\pm 10\%$ steady error

Roll Left




Roll Right



$$\begin{aligned} e_x &= x - x_{ref} \\ \ddot{e}_x &\approx T\phi \\ &= Tk_P e_x \end{aligned}$$

X-acceleration (a_x)



CHALLENGES -
IMPLEMENTATION
(FREQUENCY,
TF+ROS, OSCILLATIONS)

CHALLENGES - IMPLEMENTATION

- Frequency of image recording
 - Video over WiFi @ 30 fps --> sampled down to 5 Hz
 - Neural Net detection @ average 0.12 seconds ~ maximum 8.33 Hz frequency
- Tensorflow + ROS binding not perfect
- Roll command oscillations - tracking

NON DRONE VIDEOS – LAB
TESTING

+

REAL TIME DRONE VIDEO
– OUTDOOR FLYING

Don't Move

Box Center Location: 45.8661109209 %

Don't Move

Start

Stop

Pause

Play

Send

00:00:03 | 25 | 199 | A | 180x264 | avi | Pulse 25

AR.Drone Video Feed



person: 90%

LIMITATIONS & FUTURE WORK

LIMITATIONS & FUTURE WORK

Limitations:

- Drone is very sensitive to wind
- Loss of length reference
 - Using only 1 camera reduces everything to image plane
 - Area in image plan is highly attitude dependent

Future Work:

- Implement pitch controller (move forward & backward)
- Implement yaw controller based on orientation
- Yaw – Roll selection algorithm
- Implementation directly onto quadcopter

REFERENCES

REFERENCES

U.S. Patent No. 10,240,930. (2019). Washington, DC: U.S. Patent and Trademark Office.

K. K. Lekkala and V. K. Mittal, "Simultaneous aerial vehicle localization and Human tracking," 2016 IEEE Region 10 Conference (TENCON), Singapore, 2016, pp. 379-383. doi: 10.1109/TENCON.2016.7848025


<https://medium.com/nanonets/how-i-built-a-self-flying-drone-to-track-people-in-under-50-lines-of-code-7485de7f828e>

W. Liu, D. Anguelov, D. Erhan, C. Szegedy, S. Reed, C.-Y. Fu, A.C. Berg, Ssd: Single shot multibox detector, 2015

A. Howard, M. Zhu, B. Chen, D. Kalenichenko, W. Wang, T. Weyand, M. Andreetto, and H. Adam. Mobilenets: Efficient convolutional neural networks for mobile vision applications. arXiv preprint arXiv:1704.04861, 2017.

DEMONSTRATION


Thank
You



Questions



Reward
Function?



IT DEPENDS!